

BigMart Sales Prediction

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
In [3]: df = pd.read_csv('train.csv')
```

```
In [4]: df.head()
```

Out[4]:

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Esta
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8092	OUT049	
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2692	OUT018	
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6180	OUT049	
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	OUT010	
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614	OUT013	

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8523 entries, 0 to 8522
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                        8523 non-null   object
1   Item_Weight                           7060 non-null   float64
2   Item_Fat_Content                       8523 non-null   object
3   Item_Visibility                       8523 non-null   float64
4   Item_Type                             8523 non-null   object
5   Item_MRP                              8523 non-null   float64
6   Outlet_Identifier                      8523 non-null   object
7   Outlet_Establishment_Year             8523 non-null   int64
8   Outlet_Size                           6113 non-null   object
9   Outlet_Location_Type                  8523 non-null   object
10  Outlet_Type                           8523 non-null   object
11  Item_Outlet_Sales                     8523 non-null   float64
dtypes: float64(4), int64(1), object(7)
memory usage: 799.2+ KB
```

```
In [6]: df.describe()
```

```
Out[6]:
```

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	7060.000000	8523.000000	8523.000000	8523.000000	8523.000000
mean	12.857645	0.066132	140.992782	1997.831867	2181.288914
std	4.643456	0.051598	62.275067	8.371760	1706.499616
min	4.555000	0.000000	31.290000	1985.000000	33.290000
25%	8.773750	0.026989	93.826500	1987.000000	834.247400
50%	12.600000	0.053931	143.012800	1999.000000	1794.331000
75%	16.850000	0.094585	185.643700	2004.000000	3101.296400
max	21.350000	0.328391	266.888400	2009.000000	13086.964800

Dealing with Null Values

```
In [7]: df.shape
```

```
Out[7]: (8523, 12)
```

```
In [8]: df.isna().sum()
```

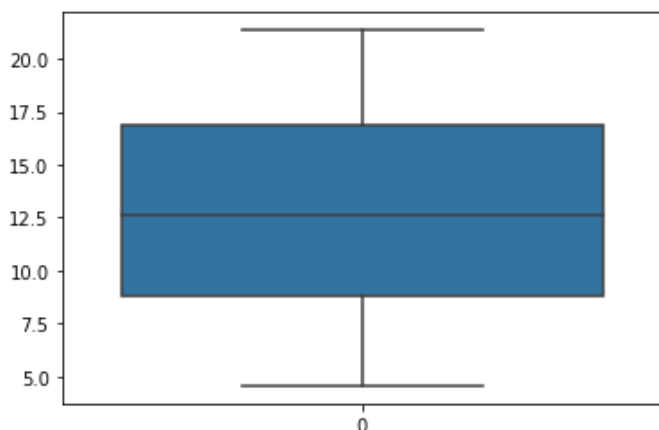
```
Out[8]: Item_Identifier      0
Item_Weight      1463
Item_Fat_Content      0
Item_Visibility      0
Item_Type          0
Item_MRP           0
Outlet_Identifier      0
Outlet_Establishment_Year  0
Outlet_Size      2410
Outlet_Location_Type      0
Outlet_Type           0
Item_Outlet_Sales      0
dtype: int64
```

```
In [9]: df['Item_Weight'].describe()
```

```
Out[9]: count      7060.000000
mean         12.857645
std           4.643456
min           4.555000
25%           8.773750
50%          12.600000
75%          16.850000
max           21.350000
Name: Item_Weight, dtype: float64
```

```
In [10]: sns.boxplot(df.Item_Weight)
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1ed1e16ba08>
```



- It seems there are no outliers, hence it is safe to impute with the mean.

```
In [11]: df.Item_Weight.fillna(df.Item_Weight.mean(),inplace=True)
```

```
In [12]: # Imputing Outlet size with mode.  
df.Outlet_Size.fillna(df.Outlet_Size.mode()[0],inplace=True)
```

```
In [13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 8523 entries, 0 to 8522  
Data columns (total 12 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   Item_Identifier                       8523 non-null   object  
1   Item_Weight                           8523 non-null   float64  
2   Item_Fat_Content                       8523 non-null   object  
3   Item_Visibility                       8523 non-null   float64  
4   Item_Type                             8523 non-null   object  
5   Item_MRP                              8523 non-null   float64  
6   Outlet_Identifier                     8523 non-null   object  
7   Outlet_Establishment_Year             8523 non-null   int64  
8   Outlet_Size                           8523 non-null   object  
9   Outlet_Location_Type                  8523 non-null   object  
10  Outlet_Type                           8523 non-null   object  
11  Item_Outlet_Sales                     8523 non-null   float64  
dtypes: float64(4), int64(1), object(7)  
memory usage: 799.2+ KB
```

Data Visualization of Numeric Columns

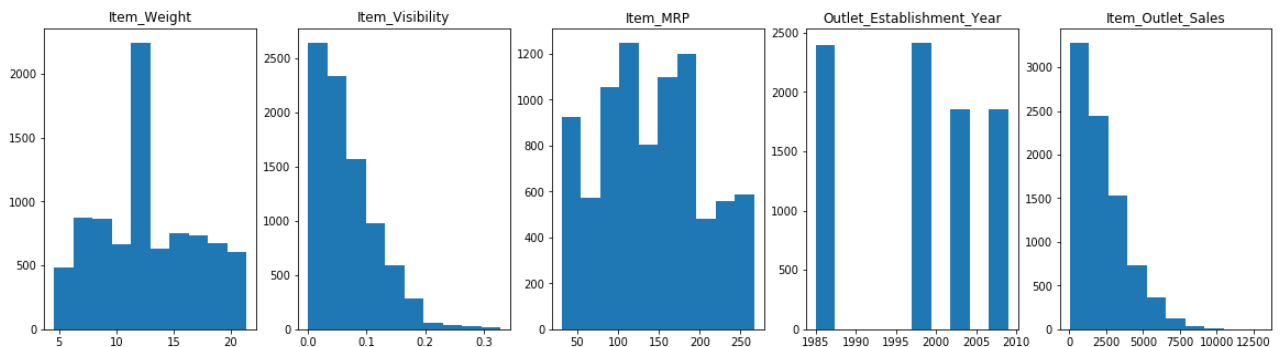
```
In [14]: nums = df.select_dtypes(include=['float64', 'int64']).columns.to_numpy()  
nums
```

```
Out[14]: array(['Item_Weight', 'Item_Visibility', 'Item_MRP',  
              'Outlet_Establishment_Year', 'Item_Outlet_Sales'], dtype=object)
```

```
In [15]: # Checking the distribution of the numeric Columns
```

```
fig, ax = plt.subplots(1, 5, figsize=(20, 5))
```

```
for i, col in enumerate(nums):  
    ax[i].hist(df[col])  
    ax[i].set_title(col)
```

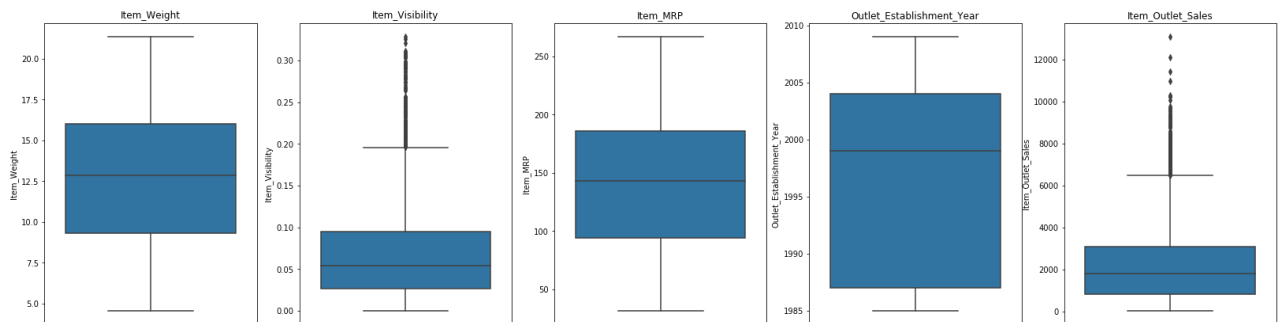


- It seems that the Item Visibility and Item outlet sales columns are right skewed!
- We can apply Data transformations to make them Normally distributed.

```
In [16]: # Checking for outliers in the numeric columns
```

```
fig, ax = plt.subplots(1, 5, figsize=(28, 7))
```

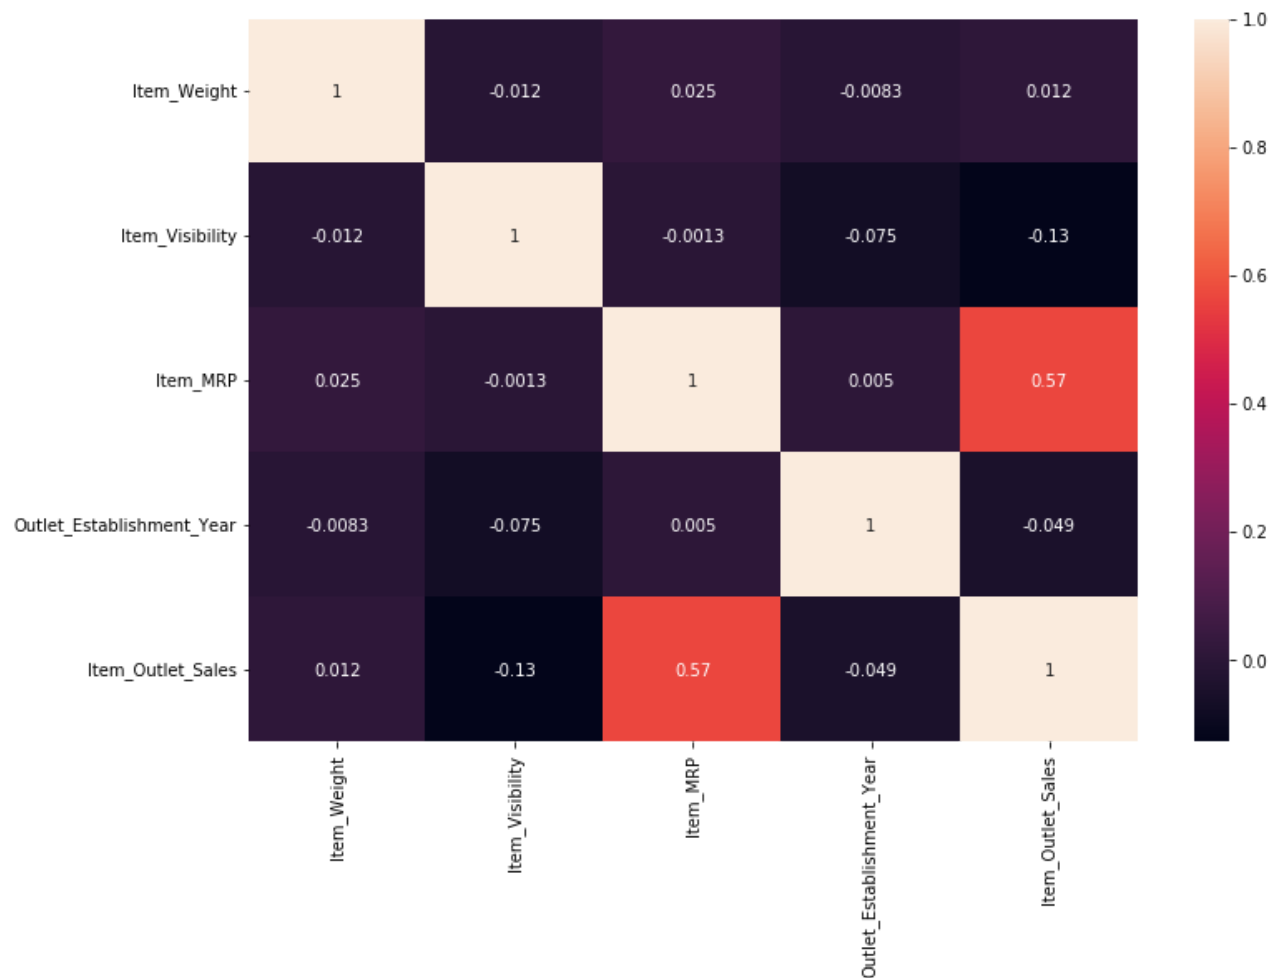
```
for i, col in enumerate(nums):  
    sns.boxplot(data=df, y=col, ax=ax[i])  
    ax[i].set_title(col)
```



- It seems that the Item Visibility and Item outlet sales columns have some outliers which need to be handled!

```
In [17]: # Observing the relation between the numeric Columns
```

```
plt.figure(figsize=(12,8))
sns.heatmap(df[nums].corr(),annot=True)
plt.show()
```



Feature Engineering of Numeric Columns

```
In [18]: def remove_outliers(col):
          q1 = df[col].quantile(0.25)
          q3 = df[col].quantile(0.75)
          iqr = q3 - q1

          upper_limit = q3 + 1.5*iqr
          count_outliers = df[df[col] > upper_limit].shape[0]

          print(f'{count_outliers} outliers out of {df.shape[0]}. i.e; {count_outliers/df.shape[0]*100}%')

          return df[df[col] < upper_limit]
```

```
In [19]: df = remove_outliers('Item_Visibility')
```

144 outliers out of 8523. i.e; 1.69% of outliers.

```
In [20]: df = remove_outliers('Item_Outlet_Sales')
```

186 outliers out of 8379. i.e; 2.22% of outliers.

```
In [21]: df.head()
```

Out[21]:

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Esta
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8092	OUT049	
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2692	OUT018	
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6180	OUT049	
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	OUT010	
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614	OUT013	

```

In [22]: fig, ax = plt.subplots(2,2,figsize=(12,12))


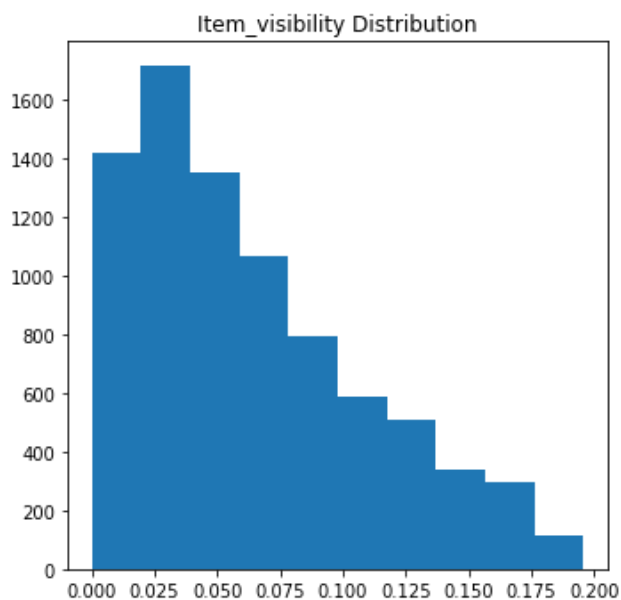

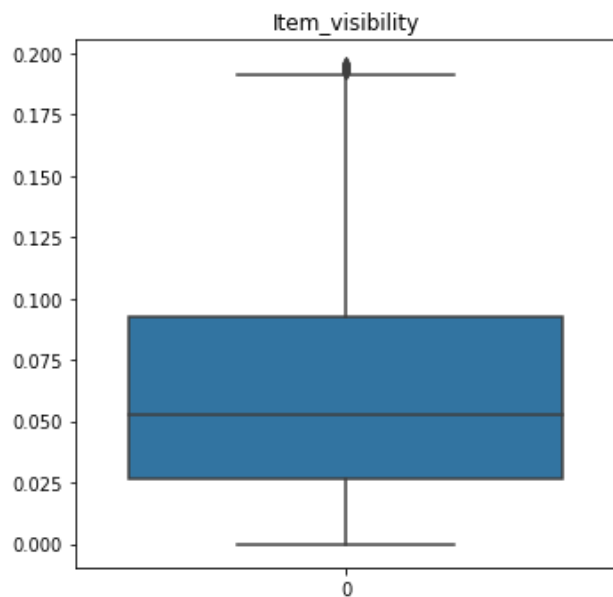
sns.boxplot(df.Item_Visibility,ax=ax[0,0])
ax[0,0].set_title('Item_visibility')

sns.boxplot(df.Item_Outlet_Sales,ax=ax[0,1])
ax[0,1].set_title('Item_Outlet_Sales')

ax[1,0].hist(df.Item_Visibility)
ax[1,0].set_title('Item_visibility Distribution')

ax[1,1].hist(df.Item_Outlet_Sales)
ax[1,1].set_title('Item_Outlet_Sales Distribution')
plt.show()

```



```
In [23]: log_item = np.log(df.Item_Visibility)
```

```
In [24]: # After Log, it is having negative values, hence to check the distribution, removing outliers.  
q1 = log_item.quantile(0.25)  
q3 = log_item.quantile(0.75)  
iqr = q3-q1  
  
l1 = iqr + 1.5*q1
```



```

In [25]: fig, ax1 = plt.subplots(2,2,figsize=(12,12))

ax1[0,0].hist(log_item[log_item>1])
ax1[0,0].set_title('Item_visibility Log Distribution')

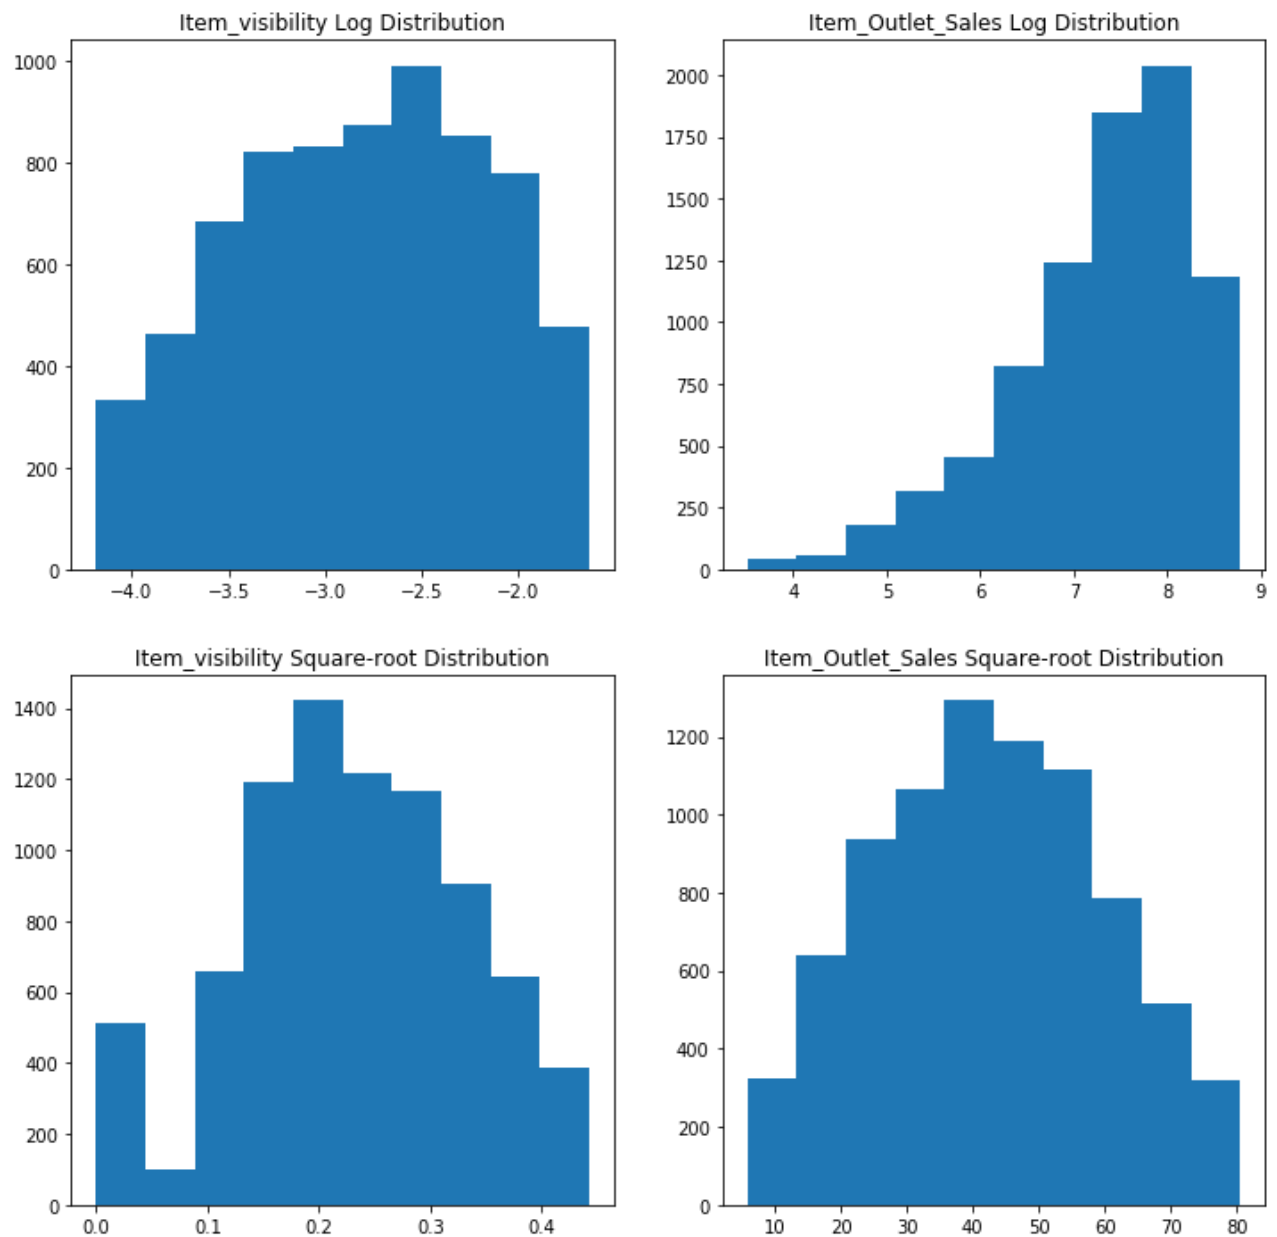
ax1[0,1].hist(np.log(df.Item_Outlet_Sales))
ax1[0,1].set_title('Item_Outlet_Sales Log Distribution')

ax1[1,0].hist(np.sqrt(df.Item_Visibility))
ax1[1,0].set_title('Item_visibility Square-root Distribution')

ax1[1,1].hist(np.sqrt(df.Item_Outlet_Sales))
ax1[1,1].set_title('Item_Outlet_Sales Square-root Distribution')

plt.show()

```




- Comparatively, we can say that applying square root have made them normally distributed. And hence, we can transform them with square root.

```
In [26]: df.Item_Visibility = np.sqrt(df.Item_Visibility)
df.Item_Outlet_Sales = np.sqrt(df.Item_Outlet_Sales)
```

```
In [27]: # Outlet establishment year will be much useful, if we can convert that to the no of years since
df['nof_years'] = 2023 - df.Outlet_Establishment_Year
```

```
In [28]: df.nof_years.value_counts().plot(kind='bar')
plt.show()
```



```
In [29]: df.drop('Outlet_Establishment_Year', axis=1, inplace=True)
```

```
In [30]: df.head()
```

```
Out[30]:
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Size
0	FDA15	9.30	Low Fat	0.126678	Dairy	249.8092	OUT049	Medium
1	DRC01	5.92	Regular	0.138846	Soft Drinks	48.2692	OUT018	Medium
2	FDN15	17.50	Low Fat	0.129461	Meat	141.6180	OUT049	Medium
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	OUT010	Medium
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614	OUT013	High

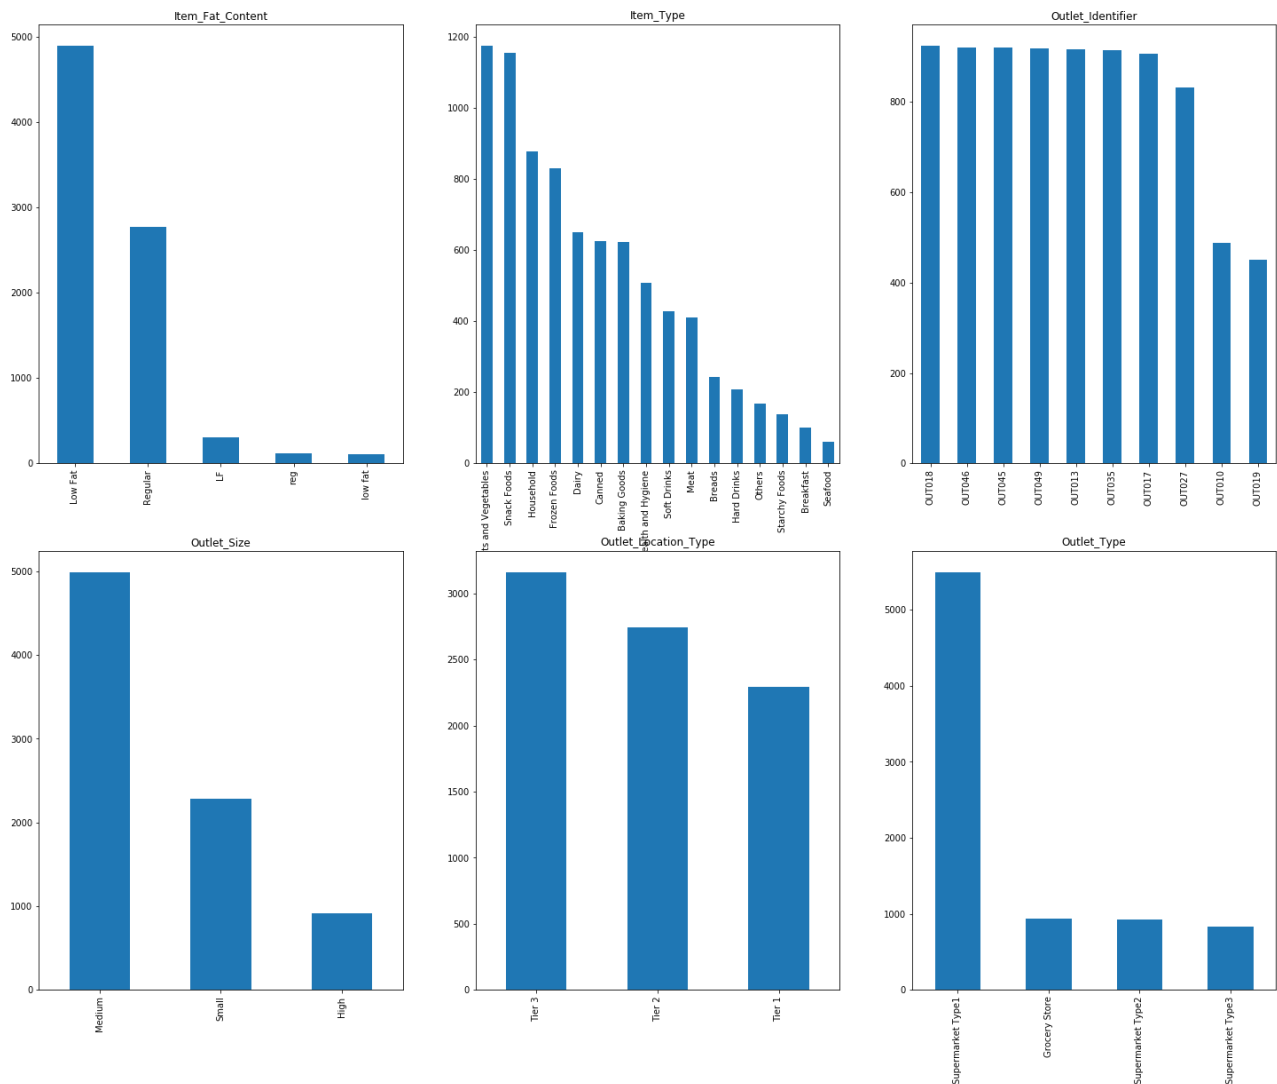
Data Visualization of Categorical Columns

```
In [31]: cat_cols = df.select_dtypes(include='object').columns.to_numpy()
```

```
In [32]: cat_cols
```

```
Out[32]: array(['Item_Identifier', 'Item_Fat_Content', 'Item_Type',  
               'Outlet_Identifier', 'Outlet_Size', 'Outlet_Location_Type',  
               'Outlet_Type'], dtype=object)
```

```
In [33]: fig, ax = plt.subplots(2,3,figsize=(25,20))  
a,b=0,0  
for i,col in enumerate(cat_cols[1:]):  
    if a!=1:  
        df[col].value_counts().plot(kind='bar',ax=ax[a,b])  
        ax[a,b].set_title(col)  
        if b==2: a,b=1,0  
        else: b+=1  
    else:  
        df[col].value_counts().plot(kind='bar',ax=ax[a,b])  
        ax[a,b].set_title(col)  
        b+=1
```



- So it seems, Low fat is recorded in 3 different names, and regular is recorded in two diff names.
- Outlet_Size, Item_Fat_Content, Outlet_Type, Outlet_Location_Type seems to be ordinal, and hence they can be label encoded.

- Item_Type, Outlet_Identifier can be one hot encoded.

Feature Engineering Categorical Columns

```
In [34]: df.select_dtypes(include='object').head()
```

```
Out[34]:
```

	Item_Identifier	Item_Fat_Content	Item_Type	Outlet_Identifier	Outlet_Size	Outlet_Location_Type	Outlet_Type
0	FDA15	Low Fat	Dairy	OUT049	Medium	Tier 1	Supermarket Type1
1	DRC01	Regular	Soft Drinks	OUT018	Medium	Tier 3	Supermarket Type2
2	FDN15	Low Fat	Meat	OUT049	Medium	Tier 1	Supermarket Type1
3	FDX07	Regular	Fruits and Vegetables	OUT010	Medium	Tier 3	Grocery Store
4	NCD19	Low Fat	Household	OUT013	High	Tier 3	Supermarket Type1

- If we see the item identifier, the first characters of it are basically the categories, so we can categorize them in that way.

```
In [35]: df['Item_Categories'] = df['Item_Identifier'].str[0:2]
df.drop('Item_Identifier',axis=1,inplace=True)
```

```
In [36]: df.Item_Fat_Content.unique()
```

```
Out[36]: array(['Low Fat', 'Regular', 'low fat', 'LF', 'reg'], dtype=object)
```

- Indeed they seem to repeated due to human error.

```
In [37]: df['Item_Fat_Content'].replace(['low fat','LF'],'Low Fat',inplace=True)
df['Item_Fat_Content'].replace('reg','Regular',inplace=True)
```

```
In [38]: df.Item_Fat_Content.value_counts(dropna=False)
```

```
Out[38]: Low Fat    5309
Regular    2884
Name: Item_Fat_Content, dtype: int64
```

```
In [39]: # Label Encoding outlet size column
df.Item_Fat_Content = df.Item_Fat_Content.map({'Regular': 0,'Low Fat': 1}).astype(int)
```

```
In [40]: df.select_dtypes(include='object')['Item_Type'].unique()
```

```
Out[40]: array(['Dairy', 'Soft Drinks', 'Meat', 'Fruits and Vegetables',
                'Household', 'Baking Goods', 'Snack Foods', 'Frozen Foods',
                'Breakfast', 'Health and Hygiene', 'Hard Drinks', 'Canned',
                'Breads', 'Starchy Foods', 'Others', 'Seafood'], dtype=object)
```

```
In [41]: # Label Encoding outlet size column
df.Outlet_Size = df.Outlet_Size.map({'Small': 1, 'Medium': 2, 'High': 3}).astype(int)
```

```
In [42]: df.head()
```

Out[42]:

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Size	Outlet_Locatio
0	9.30	1	0.126678	Dairy	249.8092	OUT049	2	
1	5.92	0	0.138846	Soft Drinks	48.2692	OUT018	2	
2	17.50	1	0.129461	Meat	141.6180	OUT049	2	
3	19.20	0	0.000000	Fruits and Vegetables	182.0950	OUT010	2	
4	8.93	1	0.000000	Household	53.8614	OUT013	3	

```
In [43]: df.select_dtypes(include='object')
```

Out[43]:

	Item_Type	Outlet_Identifier	Outlet_Location_Type	Outlet_Type	Item_Categories
0	Dairy	OUT049	Tier 1	Supermarket Type1	FD
1	Soft Drinks	OUT018	Tier 3	Supermarket Type2	DR
2	Meat	OUT049	Tier 1	Supermarket Type1	FD
3	Fruits and Vegetables	OUT010	Tier 3	Grocery Store	FD
4	Household	OUT013	Tier 3	Supermarket Type1	NC
...
8518	Snack Foods	OUT013	Tier 3	Supermarket Type1	FD
8519	Baking Goods	OUT045	Tier 2	Supermarket Type1	FD
8520	Health and Hygiene	OUT035	Tier 2	Supermarket Type1	NC
8521	Snack Foods	OUT018	Tier 3	Supermarket Type2	FD
8522	Soft Drinks	OUT046	Tier 1	Supermarket Type1	DR

8193 rows × 5 columns

```
In [44]: # Label Encoding Outlet Location type and outlet type
from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
df['Outlet_Type'] = encoder.fit_transform(df['Outlet_Type'])
df['Outlet_Location_Type'] = encoder.fit_transform(df['Outlet_Location_Type'])
```

```
In [45]: df.head()
```

Out[45]:

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Size	Outlet_Location
0	9.30	1	0.126678	Dairy	249.8092	OUT049	2	
1	5.92	0	0.138846	Soft Drinks	48.2692	OUT018	2	
2	17.50	1	0.129461	Meat	141.6180	OUT049	2	
3	19.20	0	0.000000	Fruits and Vegetables	182.0950	OUT010	2	
4	8.93	1	0.000000	Household	53.8614	OUT013	3	

```
In [46]: df.select_dtypes(include='object')
```

Out[46]:

	Item_Type	Outlet_Identifier	Item_Categories
0	Dairy	OUT049	FD
1	Soft Drinks	OUT018	DR
2	Meat	OUT049	FD
3	Fruits and Vegetables	OUT010	FD
4	Household	OUT013	NC
...
8518	Snack Foods	OUT013	FD
8519	Baking Goods	OUT045	FD
8520	Health and Hygiene	OUT035	NC
8521	Snack Foods	OUT018	FD
8522	Soft Drinks	OUT046	DR

8193 rows × 3 columns

- Now we have to deal with other cat columns as Item_Type,Outlet_Identifier,Item_categories

```
In [47]: # onehot encoding for the rest

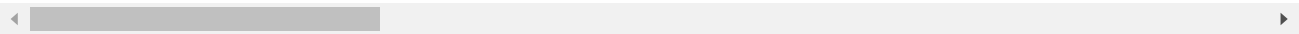
df = pd.get_dummies(df, columns=['Item_Type', 'Item_Categories', 'Outlet_Identifier'])
```

```
In [48]: df.head()
```

Out[48]:

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_MRP	Outlet_Size	Outlet_Location_Type	Outlet_Type	Item_O
0	9.30	1	0.126678	249.8092	2	0	1	
1	5.92	0	0.138846	48.2692	2	2	2	
2	17.50	1	0.129461	141.6180	2	0	1	
3	19.20	0	0.000000	182.0950	2	2	0	
4	8.93	1	0.000000	53.8614	3	2	1	

5 rows × 38 columns



```
In [49]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8193 entries, 0 to 8522
Data columns (total 38 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Item_Weight                               8193 non-null   float64
1   Item_Fat_Content                           8193 non-null   int32
2   Item_Visibility                           8193 non-null   float64
3   Item_MRP                                   8193 non-null   float64
4   Outlet_Size                               8193 non-null   int32
5   Outlet_Location_Type                       8193 non-null   int32
6   Outlet_Type                               8193 non-null   int32
7   Item_Outlet_Sales                         8193 non-null   float64
8   nof_years                                8193 non-null   int64
9   Item_Type_Baking Goods                    8193 non-null   uint8
10  Item_Type_Breads                           8193 non-null   uint8
11  Item_Type_Breakfast                        8193 non-null   uint8
12  Item_Type_Canned                          8193 non-null   uint8
13  Item_Type_Dairy                           8193 non-null   uint8
14  Item_Type_Frozen Foods                    8193 non-null   uint8
15  Item_Type_Fruits and Vegetables           8193 non-null   uint8
16  Item_Type_Hard Drinks                     8193 non-null   uint8
17  Item_Type_Health and Hygiene              8193 non-null   uint8
18  Item_Type_Household                       8193 non-null   uint8
19  Item_Type_Meat                            8193 non-null   uint8
20  Item_Type_Others                           8193 non-null   uint8
21  Item_Type_Seafood                         8193 non-null   uint8
22  Item_Type_Snack Foods                     8193 non-null   uint8
23  Item_Type_Soft Drinks                     8193 non-null   uint8
24  Item_Type_Starchy Foods                   8193 non-null   uint8
25  Item_Categories_DR                        8193 non-null   uint8
26  Item_Categories_FD                        8193 non-null   uint8
27  Item_Categories_NC                        8193 non-null   uint8
28  Outlet_Identifier_OUT010                  8193 non-null   uint8
29  Outlet_Identifier_OUT013                  8193 non-null   uint8
30  Outlet_Identifier_OUT017                  8193 non-null   uint8
31  Outlet_Identifier_OUT018                  8193 non-null   uint8
32  Outlet_Identifier_OUT019                  8193 non-null   uint8
33  Outlet_Identifier_OUT027                  8193 non-null   uint8
34  Outlet_Identifier_OUT035                  8193 non-null   uint8
35  Outlet_Identifier_OUT045                  8193 non-null   uint8
36  Outlet_Identifier_OUT046                  8193 non-null   uint8
37  Outlet_Identifier_OUT049                  8193 non-null   uint8
dtypes: float64(4), int32(4), int64(1), uint8(29)
memory usage: 1.0 MB
```

Model Building

```
In [50]: from sklearn.model_selection import train_test_split
         from sklearn.metrics import r2_score
```

```
In [51]: X = df.drop('Item_Outlet_Sales', axis=1)
         y = df['Item_Outlet_Sales']
```



```
In [52]: X_train,X_test,y_train,y_test = train_test_split(X,y)
```

```
In [53]: def scores(model):  
    train_pred = model.predict(X_train)  
    print(f'Score on Training dataset = {r2_score(y_train,train_pred)}')  
    test_pred = model.predict(X_test)  
    print(f'Score on Training dataset = {r2_score(y_test,test_pred)}')
```

Simple Linear Regression

```
In [54]: from sklearn.linear_model import LinearRegression  
  
lin_reg = LinearRegression()  
lin_reg.fit(X_train, y_train)  
scores(lin_reg)
```

Score on Training dataset = 0.6383938275642018
Score on Training dataset = 0.6214603498517137

Applying Regularization Techniques

The key difference is in how they assign penalties to the coefficients:

Ridge Regression:

1. Performs L2 regularization, i.e., adds penalty equivalent to the square of the magnitude of coefficients
2. Minimization objective = LS Obj + α * (sum of square of coefficients)

Lasso Regression:

1. Performs L1 regularization, i.e., adds penalty equivalent to the absolute value of the magnitude of coefficients
 2. Minimization objective = LS Obj + α * (sum of the absolute value of coefficients)
- Here, LS Obj refers to the 'least squares objective,' i.e., the linear regression objective without regularization.

```
In [55]: from sklearn.linear_model import Ridge  
  
ridge = Ridge()  
ridge.fit(X_train,y_train)  
scores(ridge)
```

Score on Training dataset = 0.6383934083247518
Score on Training dataset = 0.6214350758753172

```
In [56]: from sklearn.linear_model import Lasso
```

```
lasso = Lasso()  
lasso.fit(X_train,y_train)  
scores(lasso)
```

Score on Training dataset = 0.4981745084448651
Score on Training dataset = 0.46019129104610856

Bossting Regressors

```
In [57]: from sklearn.ensemble import RandomForestRegressor
```

```
rfg = RandomForestRegressor()  
rfg.fit(X_train,y_train)  
scores(rfg)
```

Score on Training dataset = 0.9470717118105295
Score on Training dataset = 0.6022049030452542

- The Difference between the training score and the test score is huge, which indicatest that it is overfitting the data.
- We can get the better scores for all these algorithms, by find the right suitable paramters through Hyperparameter Tuning.

Hyperparameter Tuning the Models

Hyperparameter Tuning Ridge

```
In [58]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
```

```
# Initialize the Ridge model  
ridge = Ridge()  
  
# Define the hyperparameters to tune  
param_grid = {  
    'alpha': [0.1, 1,3,5,7,9,10],  
    'fit_intercept': [True]  
}  
  
# Perform Grid Search Cross-Validation  
grid_search = GridSearchCV(ridge, param_grid, cv=5)  
grid_search.fit(X_train, y_train)  
  
# Print the best parameters and the corresponding score  
print("Best parameters: ", grid_search.best_params_)  
print("Best score: ", grid_search.best_score_)
```

Best parameters: {'alpha': 10, 'fit_intercept': True}
Best score: 0.6351441341336047

```
In [59]: ridge_best = Ridge(alpha=10)
ridge_best.fit(X_train,y_train)
scores(ridge_best)
```

Score on Training dataset = 0.6383536983989735
Score on Training dataset = 0.6211808846712839

Hyperparameter Tuning Lasso

```
In [60]: # Create a Lasso regressor
lasso = Lasso()

# Define the hyperparameter grid
param_grid = {
    'alpha': [0.1, 1.0, 10.0, 100.0],
    'fit_intercept': [True, False],
    'normalize': [True, False],
    'max_iter': [1000, 2000, 3000]
}

# Perform Grid Search with cross-validation
grid_search = GridSearchCV(estimator=lasso, param_grid=param_grid, cv=5, verbose=2, n_jobs=-1)

# Fit the model
grid_search.fit(X_train, y_train)

# Get the best parameters
best_params = grid_search.best_params_
print("Best parameters:", best_params)
```

Fitting 5 folds for each of 48 candidates, totalling 240 fits
Best parameters: {'alpha': 0.1, 'fit_intercept': True, 'max_iter': 1000, 'normalize': False}

```
In [61]: # Use the best parameters to create a new Lasso model
best_lasso = Lasso(**best_params)
best_lasso.fit(X_train,y_train)
scores(best_lasso)
```

Score on Training dataset = 0.6346292560442294
Score on Training dataset = 0.6184364337716537

Hyperparameter Tuning RandomForest Regressor

```
In [62]: # Create a RandomForestRegressor
rf = RandomForestRegressor()

# Define the hyperparameter grid with fewer parameters
param_grid = {
    'n_estimators': [100, 300, 500],
    'max_depth': [5, 10],
    'min_samples_split': [2, 5, 10],
}

# Perform Randomized Search with cross-validation
random_search = RandomizedSearchCV(estimator=rf, param_distributions=param_grid, n_iter=100, )

# Fit the model
random_search.fit(X_train, y_train)

# Get the best parameters
best_params = random_search.best_params_

# Use the best parameters to create a new RandomForestRegressor model
best_rf = RandomForestRegressor(**best_params)

# Fit the model with the best parameters
best_rf.fit(X_train, y_train)

scores(best_rf)
```

Fitting 5 folds for each of 18 candidates, totalling 90 fits
Score on Training dataset = 0.6660287032459365
Score on Training dataset = 0.6426640021495466

```
In [63]: best_params
```

```
Out[63]: {'n_estimators': 500, 'min_samples_split': 2, 'max_depth': 5}
```

- Of All these, we can random forest gave the best score after hyperparameter tuning. So we take the random forest to be as best model Regressor to be as best model.

Saving the model

```
In [64]: import joblib

# Save the trained model
joblib.dump(best_rf, 'random_forest_regressor.pkl')
```

```
Out[64]: ['random_forest_regressor.pkl']
```